
ROCKSTARR CONTENT QUICKSTART

v0.7.3

BEFORE YOUR FIRST SESSION

- Your approved style guide in `/rockstarr-ai/00_intake/style-guide.md`.
- Your stack captured in `/rockstarr-ai/00_intake/stack.md`, including the content-cadence fields (`blogs_per_month` , `thought_leadership_per_month` , `email_newsletters_per_month` , `linkedin_newsletters_per_month` , `case_studies_per_quarter` , `records_videos`) and the `website_base_url` field that the blog lane needs for internal linking.
- At least a handful of first-party KB files processed in `/rockstarr-ai/01_knowledge_base/processed/` so the bot has voice signal to draw from.
- 60 minutes on the calendar with your Rockstarr lead for the first session — 30 minutes for the SEO strategy run, 30 minutes for the first monthly planning pass.
- Decide which month is month one. The calendar runs month-at-a-time.

First session with your Rockstarr lead (60 min)

1. We run seo-strategy first. The bot researches your website + the keyword landscape, generates 30 seed keywords and 50 long-tail variants, identifies 5 competitor gaps, and builds 4-5 topic clusters with one pillar page + 4-6 supporting blog posts each. Output: a strategy document plus a canonical backlog of 25-32 prioritized blog topics in `/02_inputs/seo/backlog.md`.
2. We confirm cadence in `stack.md`. How many blogs, thought-leadership pieces, and newsletters per month. How many case studies per quarter. Whether you record video.
3. We run ideate-topics. The bot proposes 8 to 12 ranked angles. Blog angles draw from the SEO backlog first; thought-leadership angles get an enemy field plus a diversity check; other lanes ideate from first-party material. You pick which to draft.
4. We run content-calendar. Your picks get slotted across the month. Blog and thought-leadership outlines lead their drafts. Pillar blogs calendar before their supporting posts. Newsletters anchor to your preferred weekday.

5. You approve the calendar. That is the monthly gate — once approved, per-piece drafting runs on the scheduled dates.
6. Case studies sit outside the monthly calendar and fire on a quarterly reminder when you name a customer win to document.

Your monthly rhythm (about 2 hours total)

- Day 1 of the month, 30 min. Monthly planning session — ideate, pick, calendar, approve.
- Days 4 through end-of-month, 20 to 40 min per draft. Review each draft that lands in 03_drafts/content/. Approve, edit, or send back.
- One Friday review, 30 min. Walk this week's drafts, confirm what ships next week, talk next month's pillars.
- Quarterly, 45 min. Re-run seo-strategy to refresh the backlog and run one 30-minute case-study interview per customer win.

What to type

Plain English. Type these into a Cowork task, one at a time. Your own wording works too; these are starting points.

- **Strategic planning (quarterly):** "Run the SEO strategy." The bot produces the strategy doc and refreshes the blog backlog.
- **Monthly kickoff:** "Ideate this month's content topics." Then "Build the content calendar for May." Then "Approve the calendar."
- **Blog lane:** "Outline the blog on [topic]." The bot runs WebSearch research, drafts the FAQ section, plans keyword placement and internal links. After you approve the outline, "Draft the blog on [topic]." Every external fact gets an inline [Source: URL].
- **Thought leadership:** "Outline the thought-leadership piece on [topic]." The bot forces five required fields — thesis, smart-competitor counter-argument, opening scene, quotable line, buried proprietary term. After you approve, "Draft the thought-leadership piece."
- **Newsletter:** "Draft this week's email newsletter." Pulls the month's approved pieces as the CTA.
- **Case study:** "Draft a case study on [customer name]." The bot walks the interview one question at a time.
- **Repurpose:** "Repurpose the approved [slug] piece." Emits LinkedIn post, newsletter highlight, X/Threads thread, and (if you record video) a short script.
- **Approval:** "Approve this draft." Moves the file from 03_drafts/ to 04_approved/ with a date stamp.

- **Publishing:** "Log that I published [slug] to [channel]." Records the send in 05_published/_publish.log.
- **Any question:** Just describe what you want. "How many backlog items are left?" works. So does "Which TL angles for May have rhyming enemies?"

Your Friday rhythm (30 min)

- Open any drafts currently sitting in 03_drafts/content/ and walk them with your Rockstarr lead.
- Confirm what ships next week and what gets bumped.
- Review the month's publish log so repeat topics do not creep in.

If something goes sideways

- A blog draft is missing the FAQ section. The 13-item SEO/GEO checklist would have failed it on Pass 1 — re-running the draft from the same approved outline picks the missing piece up.
- A thought-leadership piece reads like positioning instead of perspective. The TL rubric should have caught it. Check the tl_rubric_score in front-matter; if low, re-open the outline rather than re-drafting from the same fuzzy thesis.
- A draft reads like a model wrote it. Check the stop_slop_score in front-matter. Scores below 35/50 flag for review — tell your Rockstarr lead.
- The bot refuses to draft on a lane. You probably have that lane's cadence set to 0 in stack.md (or for blogs, website_base_url is missing). Update the stack, don't work around it.
- Two TL angles for the same month have rhyming enemies. ideate-topics flags this at proposal time and asks you to resolve before approving the calendar.
- The case-study interview stalled mid-way. Re-run the skill — it resumes at the first unanswered required question from the transcript.
- The SEO backlog is running thin. ideate-topics warns when fewer than 2x monthly blog cadence remains. Re-run seo-strategy to refresh.
- Anything else. Call your Rockstarr lead.

YOUR SAFETY NET

- Every draft runs the shared stop-slop skill as its mandatory final pass. AI tells get stripped before the file lands.
- Blog drafts run the 13-item SEO/GEO quality checklist as Pass 1 before stop-slop runs as Pass 2 — structural quality first, prose hygiene second.
- Thought-leadership drafts run the canonical TL rubric as Pass 1 before stop-slop runs as Pass 2 — argument quality first, prose hygiene second.
- Both long-form lanes are outline-first. Drafting refuses to run without an approved outline. Catches fuzziness before drafting wastes compute.
- Nothing promotes to 04_approved/ without your explicit approval.
- Case studies are interview-first. The bot will not accept a pre-written case-study document.
- Short-form content (LinkedIn polls, social posts) lives in the rockstarr-social plugin, not here. This plugin is long-form only as of v0.7.