
QUICKSTART

Rockstarr Reply

BEFORE YOUR FIRST REPLY

An outreach plugin installed in your workspace (today, that is rockstarr-outreach-salesnav). Reply rides on top — it does not run on its own.

Your approved style guide and ICP qualifications, both produced during onboarding. The bot has no opinion of its own about voice or who counts as a target — it reads yours.

A booking link (Calendly, GrowthAmp, etc.) OR access to share your Google Calendar — used when a hot reply lands and the bot proposes meeting times.

An email address you actively watch — that is where urgent reply notifications and the daily digest land.

30 minutes on the calendar with your Rockstarr lead to walk through your first few approvals together.

How you talk to the system

You type short, plain-English instructions into a Cowork task. No commands, no menus. The bot drafts the reply, runs an AI-slop scrub, and shows you the draft along with who the lead is, what they said, the temperature read, the proposed label, and the proposed follow-up timer. Nothing sends until you say so. You can edit, push back, or hand the thread off entirely.

Your first reply, end to end

- **An urgent email arrives.** Subject reads something like “Reply from Jane Doe needs your review.” Inside: the inbound message, the proposed reply body, and an Open in Cowork link.
- **Click Open in Cowork.** Claude Desktop opens with the draft already loaded. You did not need to navigate anywhere.
- **Read the draft, the lead context, and the temperature read.** Pick one response: “Approve and send,” “Make it shorter,” “Try a different angle,” “Draft a breakup message,” “Flag for review,” or “Skip.”

- **Editing instructions trigger a redraft, never a send.** You will see the new version. You can keep editing as many times as you want.
- **When the draft reads right, type “Send it.”** Only “send it” (or a clear equivalent like “yes send” or “go ahead and send it”) authorizes a send. Editing instructions never do — that is the gate, and it is non-negotiable.
- **The reply ships through your outreach plugin on the right channel and account.** The verbatim words you used to authorize get logged so we have a clear record of what you approved.

The three-option flow (when a non-target says yes)

LinkedIn is full of polite people who say “absolutely” to a soft opener even when they are not a buyer. The bot will NOT auto-pitch them. Instead, it shows you three options with no default pre-selected:

- **Option A — Graceful exit.** A short, warm message naming the fit mismatch in your own positioning language.
- **Option B — Let it hang.** No reply sent. Label Ignore. The thread closes naturally.
- **Option C — Throwaway question.** A light, curiosity-preserving question that keeps the relationship warm without committing to a pitch.

You pick. The bot writes. You approve.

Your daily rhythm (10 to 15 min)

- **Throughout the day.** When a reply lands, you get an urgent email. Open the deep-link, approve or edit, send it. Total time per reply: about 60 seconds.
- **Once a day, 6am.** A daily digest email lists everything still pending across every Rockstarr bot — reply drafts you have not yet acted on, blog drafts, outreach campaign specs. One pass through the digest clears your queue.
- **If the inbox is quiet.** No replies, no email, no work. The digest also stays silent on a zero-pending day.

Other things you will type

- **Show me waiting drafts:** “Show me the replies waiting for my approval.”
- **Manually flag a thread:** “Flag the Jane Doe thread — I want to handle it myself.”
- **Mark a manual booking:** “Mark the meeting with Jane Doe as booked for Thursday at 2.”
- **Refresh your warm-reply pattern:** “Re-run the warm-reply pattern interview — the closing move feels off.”

- **Update ICP rules:** “My ICP changed — let’s tighten the rules.”
- **Any question:** Describe what you want. “Why did the bot flag the Acme thread?” and “Show me the last reply I approved on the Sales Nav account.” both work.

If something goes sideways

- **The drafts feel generic.** Re-run “capture-warm-reply-pattern” with your Rockstarr lead. Specific answers in that interview produce specific drafts.
- **The bot keeps drafting for non-targets.** Tighten icp-qualifications.md. Type “my ICP changed” and walk through the interview again.
- **A reply slipped through that you would have killed.** Tell your Rockstarr lead. We re-tune the gate language and catch the misclassification.
- **You are traveling and cannot approve for a few days.** Nothing sends. The drafts queue. The next morning’s digest will list everything still pending. No damage done.
- **Anything else.** Call your Rockstarr lead.

YOUR SAFETY NET

Nothing sends on a reply thread without your explicit “send it.” Editing instructions are NEVER authorization. Silence is never consent.

The bot has no hardcoded reply pattern — every word comes from your style guide, your warm-reply pattern, and your ICP rules. If the drafts feel off, the inputs are the place to fix it.

Every draft passes a voice pass and an AI-slop scrub before it hits your queue. The scrub returns a 0 to 50 score; anything below 35 gets flagged so you see it.

The booking link is never pasted into a reply. The bot proposes specific times instead, then your outreach plugin books on the lead’s behalf once they confirm.

Every approved send is logged with the verbatim phrase you used to authorize it. There is a clear paper trail.